

# Integrated Analysis of Temporal Behavior of Component-based Distributed Real-time Embedded Systems

Pranav Srinivas Kumar and Gabor Karsai  
Vanderbilt University/ISIS

# Introduction

---

- ▶ The need for design-time schedulability analysis and verification:
    - ▶ Hard Real-time Systems must meet operational deadlines, that constrain the amount of time permitted to elapse between a stimulus provided to the system and a response generated by the system
    - ▶ Delayed responses and missed deadlines can cause catastrophic effects on the function of the system
  - ▶ Goal: Complete model-based development toolchain: modeling, analysis, synthesis, operations, maintenance
-

# Problem Statement

---

- ▶ Many of the existing schedulability analysis tools are not directly applicable to all system designs
    - ▶ Domain-specific properties such as arbitrary component interaction patterns, distributed deployment, and time-varying networks make this problematic
  - ▶ The classic thread-based concurrency model is too low-level and too generic
    - ▶ Hard to analyze and use
  - ▶ Restrictive, yet useful concurrency and component models are needed for which dedicated analysis tools *can* be developed
-

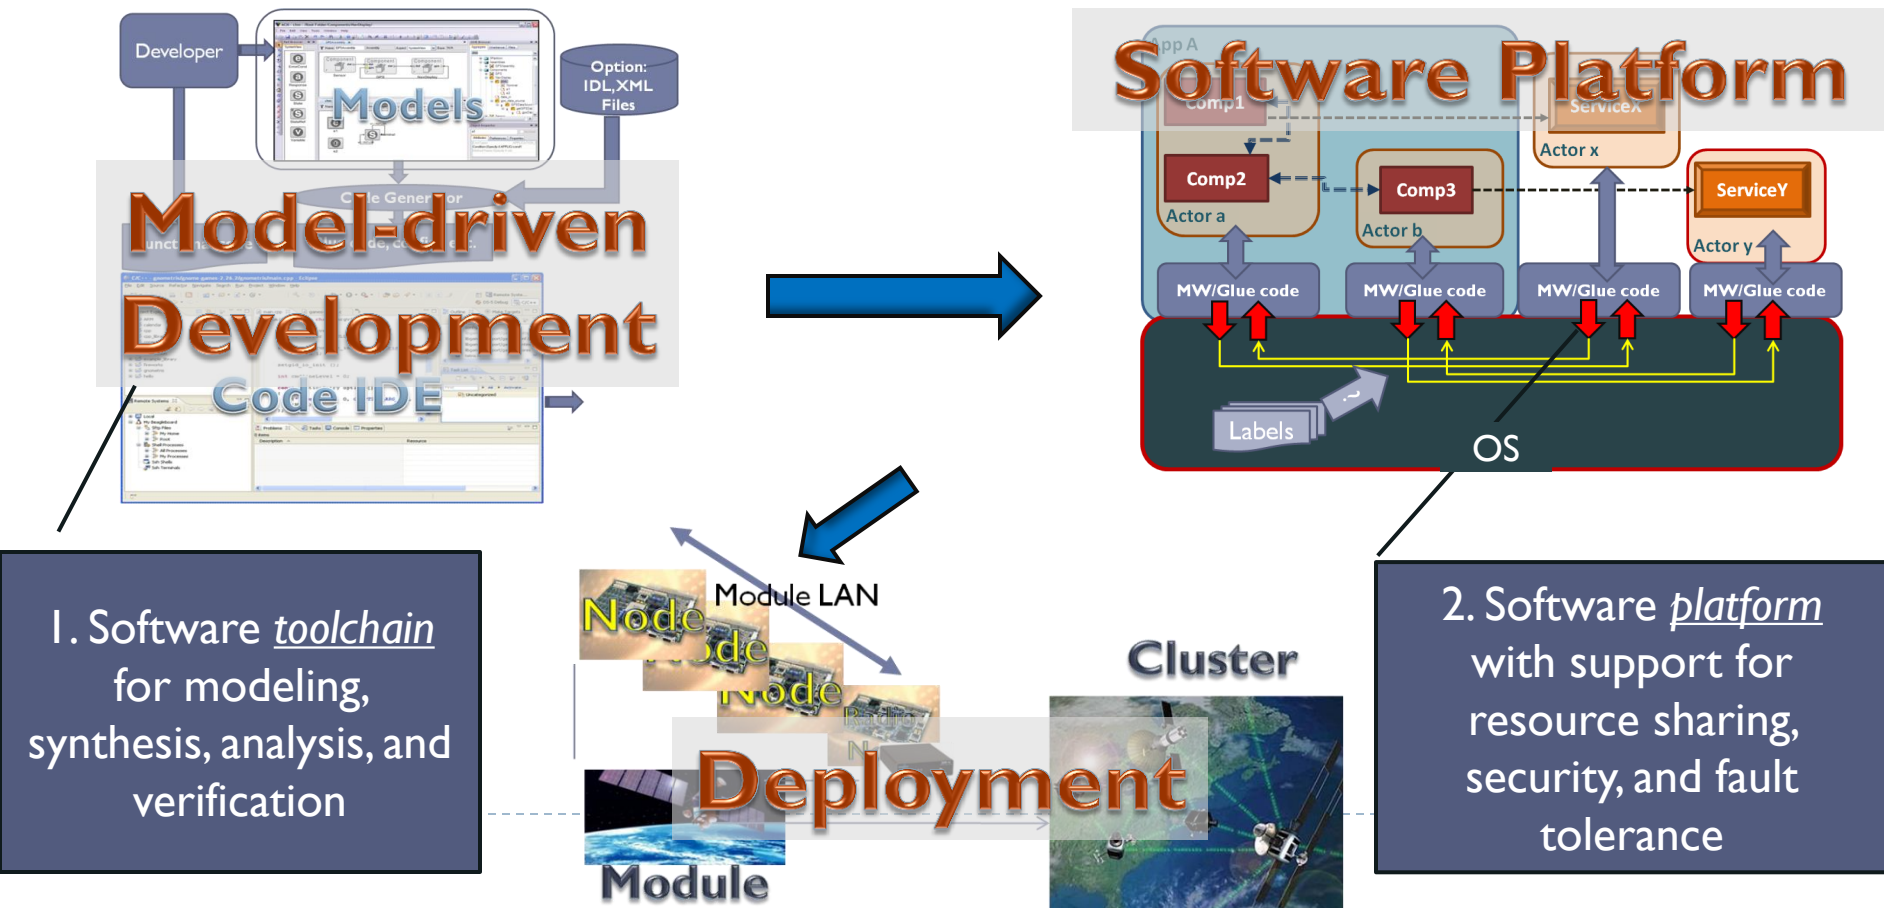
# Problem Statement

---

- ▶ Mixed-criticality Component-based applications
  - ▶ Distributed Deployment
  - ▶ Each component exposes a set of interfaces to other components and to the underlying framework
  - ▶ Different Component Interaction Patterns
  - ▶ Hierarchical Scheduling
  - ▶ Can we verify that no component operation misses its deadline?
  - ▶ Can we verify that all timing requirements are always met?
-

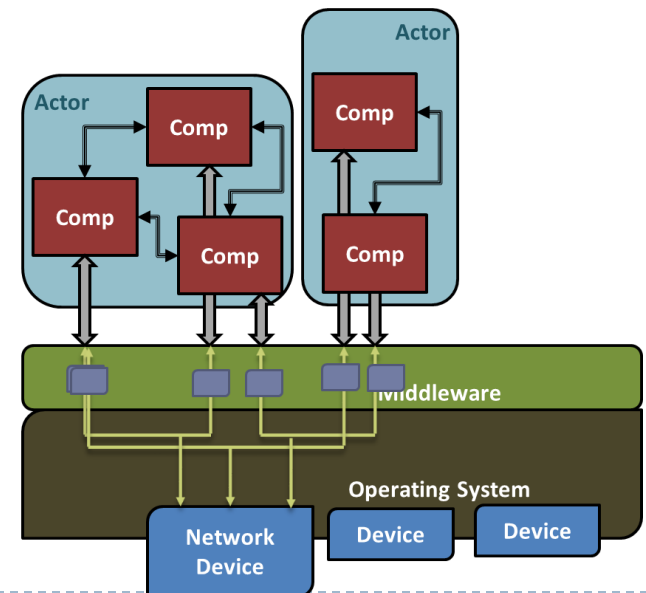
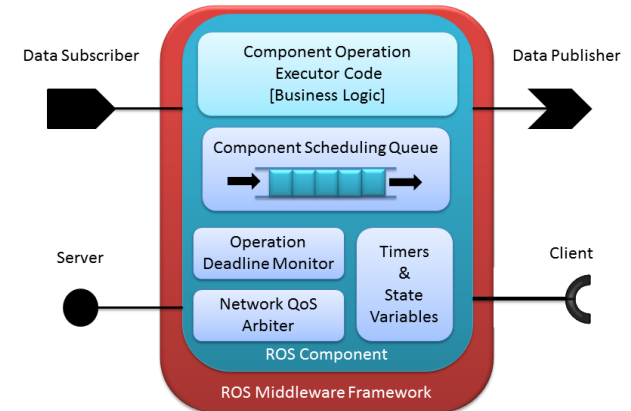
# Target Architecture: DREMS

- **Distributed REaltime Managed Systems [DREMS]** consists of:
1. A software *platform*, consisting of an OS and middleware
  2. A software *toolchain*, for modeling applications

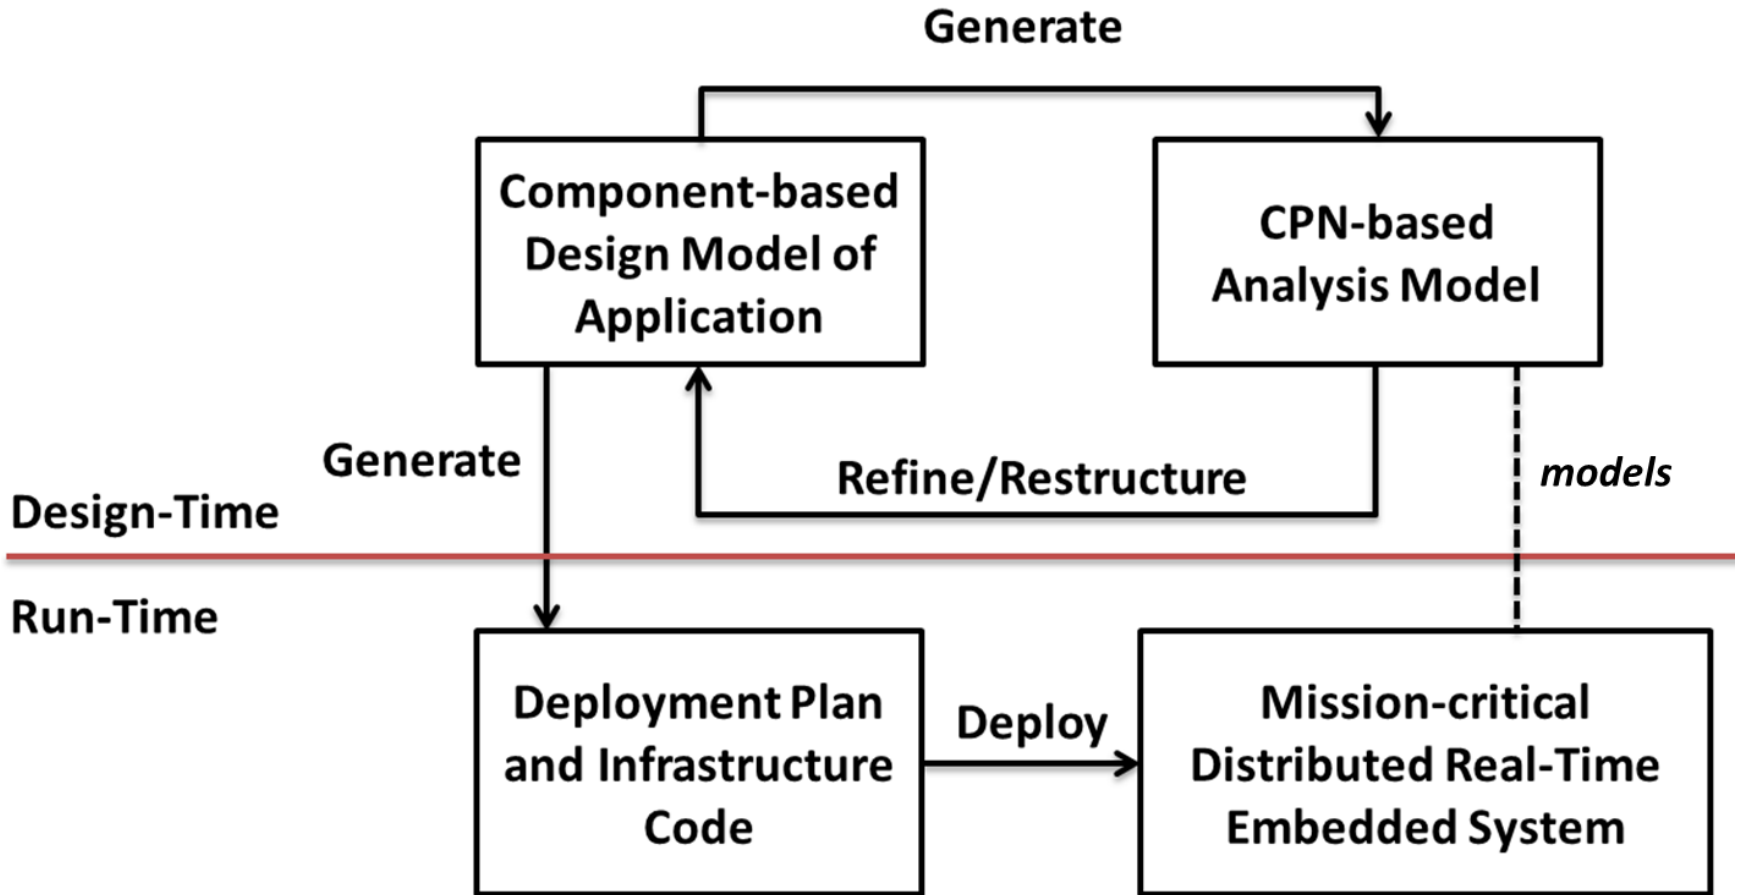


# Target Architecture: DREMS

- ▶ Developed for a class of distributed real-time embedded systems
  - ▶ Remotely managed
  - ▶ Satisfy timing requirements
- ▶ Software Infrastructure
  - ▶ Design-time: Modeling and Analysis Tools
  - ▶ Run-time: Well-defined Component Model
  - ▶ Rapid Prototyping and Code Generation Features



# Verification-driven Workflow



# Contributions

---

[Based on previous work: DREMS component model]

- ▶ A Colored Petri net-based approach to modeling and analyzing the structural behavioral properties of Component-based DRE Systems, such as DREMS
  - ▶ An approach for modeling the operational behavior of each component in an application
    - ▶ The model uses a sequence of timed steps that are executed by the operations
  - ▶ Improvements to a CPN-based modeling approach enabling better analysis performance and scalability
    - ▶ Relies on heuristics that manage time variables and state space data structures efficiently
  - ▶ Advanced state space analysis techniques applied to reduce analysis time on medium-to-large systems
-



# DREMS Background - Hierarchical Scheduling

---

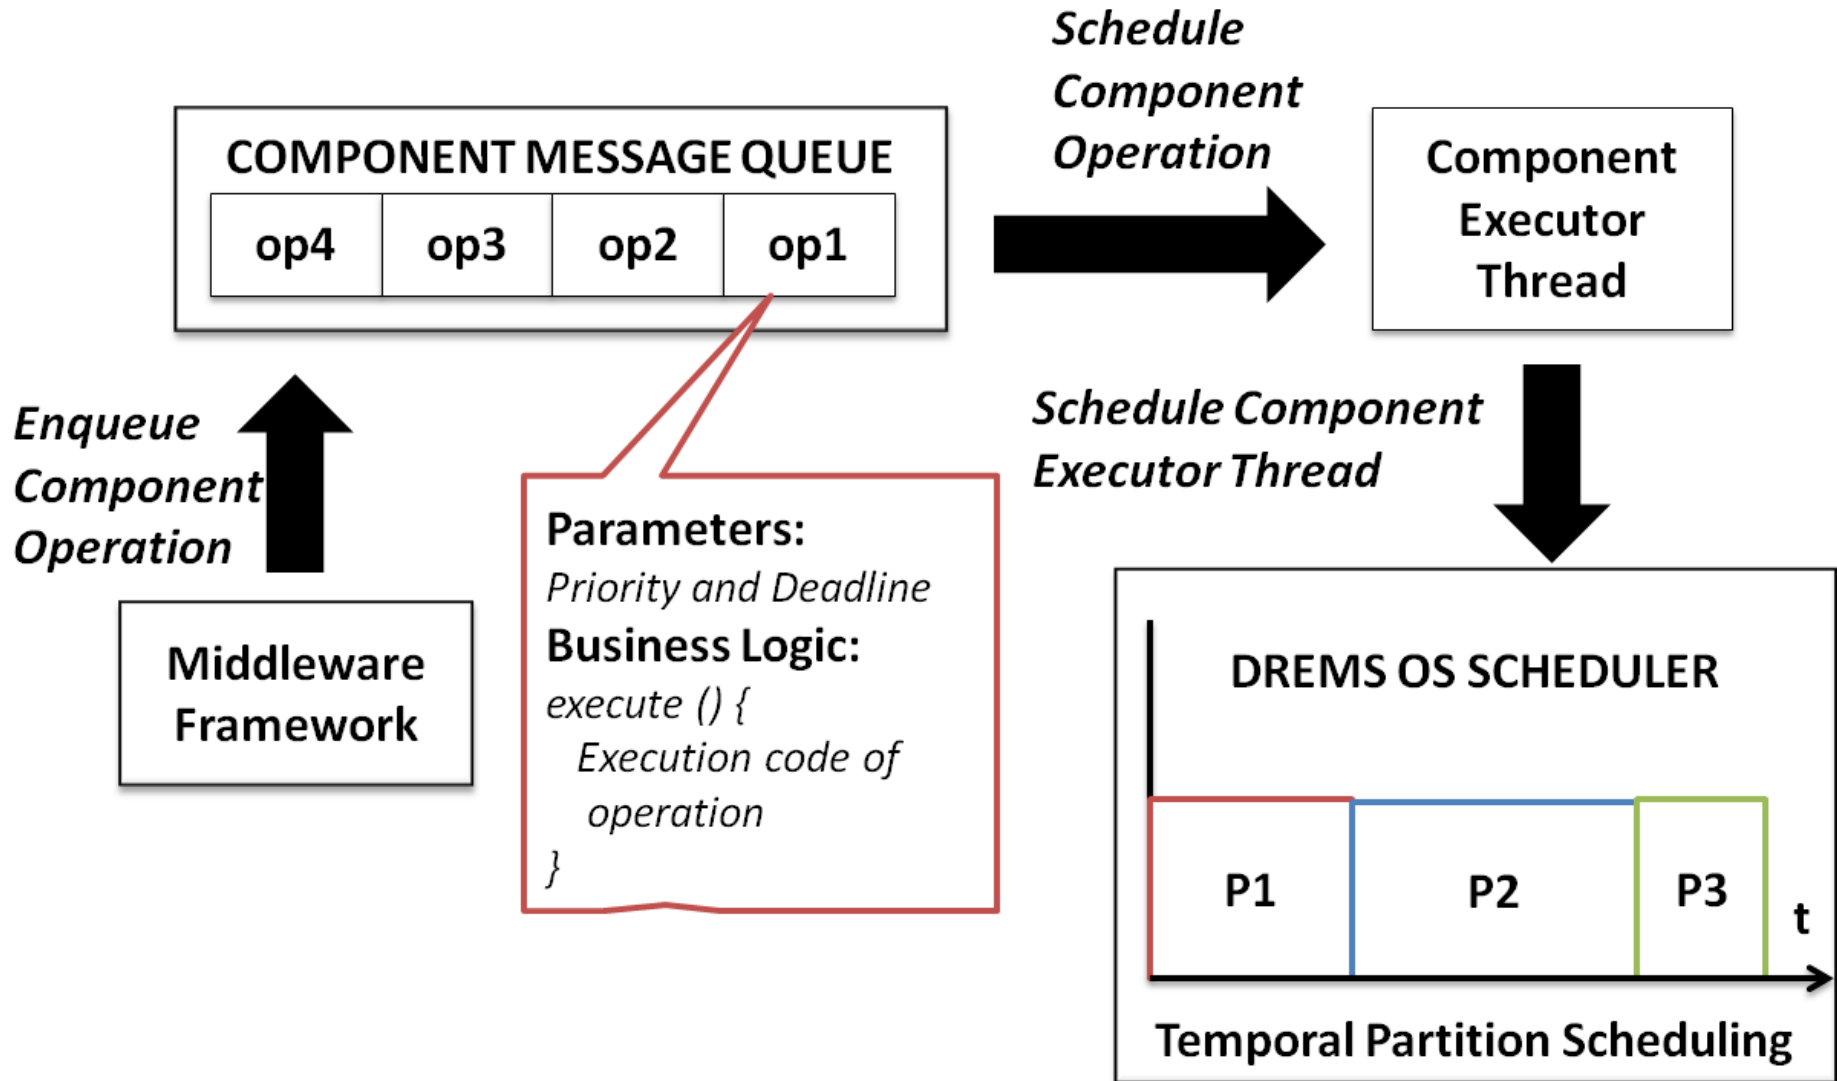
## ▶ Component level

- ▶ DREMS Applications are built by assembling *Components*
- ▶ Component interfaces expose *operations* that can be invoked
- ▶ Components interact via RMI/AMI and pub/sub interactions
- ▶ Operation requests are handled using a per-component *message queue*
- ▶ Scheduling Policy for requests: Non-preemptive Priority FIFO
  - ▶ EDF and FIFO policies are planned
- ▶ Single executor thread per component handles requests

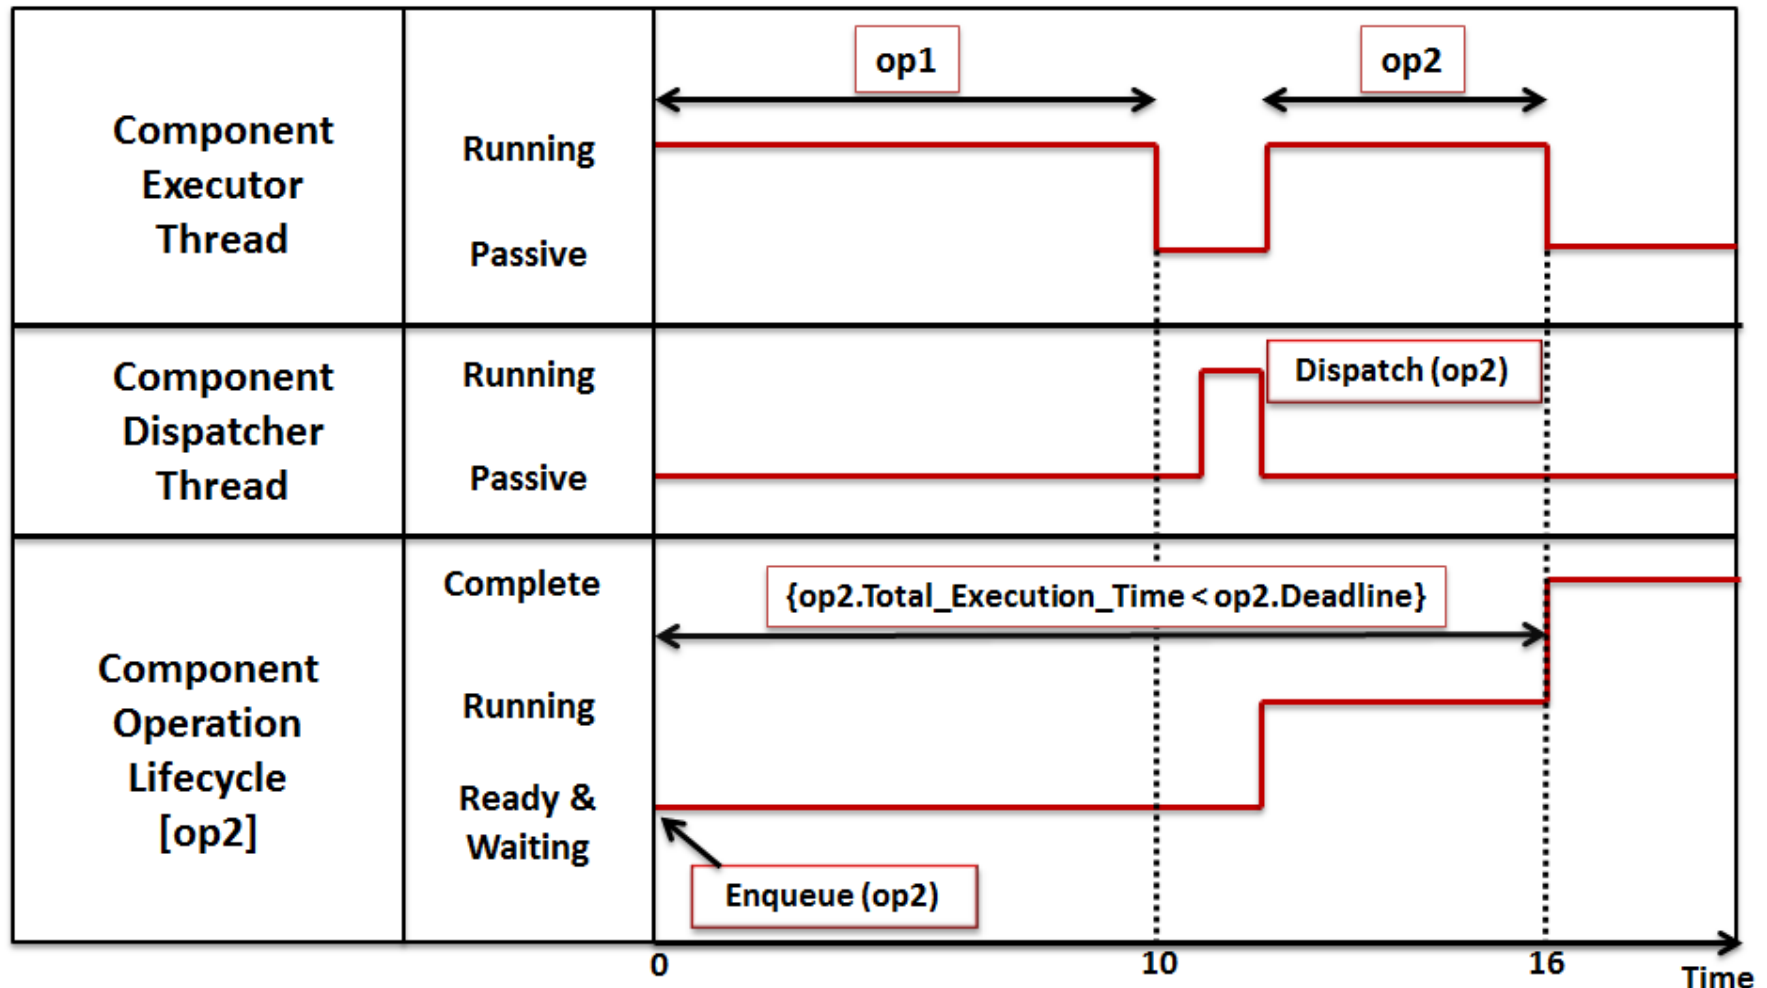
## ▶ Operating System level

- ▶ Components are grouped into processes
  - ▶ Processes are assigned to ARINC-653 style temporal partitions: fixed length, periodic intervals of the CPU's time
-

# Component Operations



# Component Operations



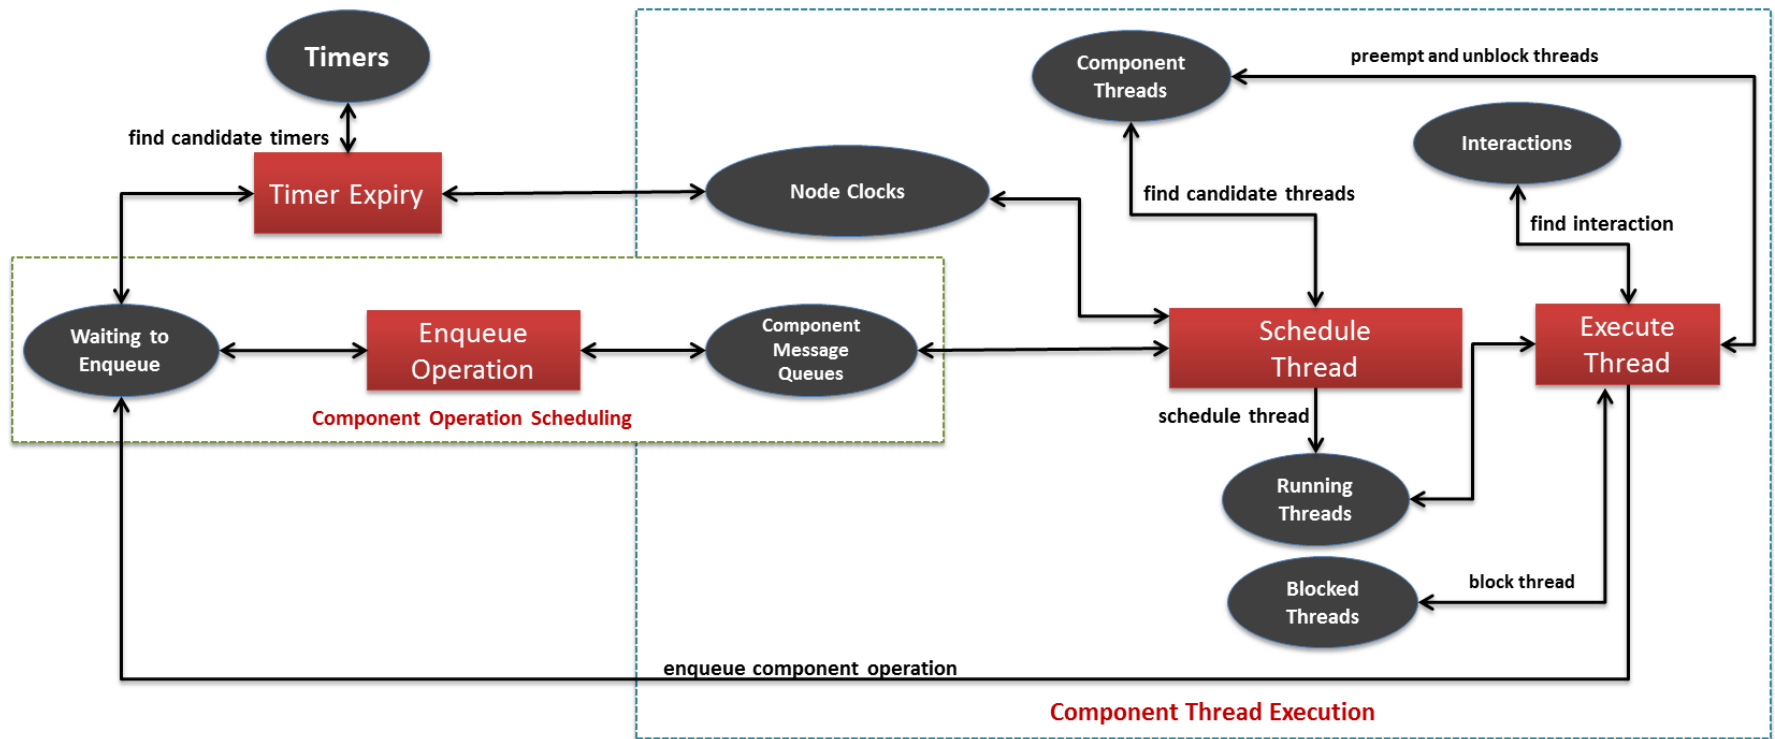
# Colored Petri Nets

---

- ▶ Graphical Modeling/Analysis Tool [[cpntools.org](http://cpntools.org)]
  - ▶ Tokens contain values of data types called *colors*
  - ▶ Powerful modeling concepts are facilitated by token colors
    - ▶ Heterogeneous data structures such as records with arbitrary number of fields
    - ▶ Tokens can be inspected, modified and manipulated by transitions and arc bindings
    - ▶ Component properties such as thread priority, port connections and real-time requirements are encoded into a single color token
-

# CPN Analysis Model

- ▶ A completed generated Colored Petri Net timing analysis model for a Domain-specific modeling language



# Modeling Temporal Behavior

---

(\* **Business Logic syntax in Extended Backus-Naur Form** \*)

```
business_logic      =   'Do', ws, operation_name, ws
                        '[' , operation_priority, operation_deadline, ']', '{', { functional_step }, '}' ;

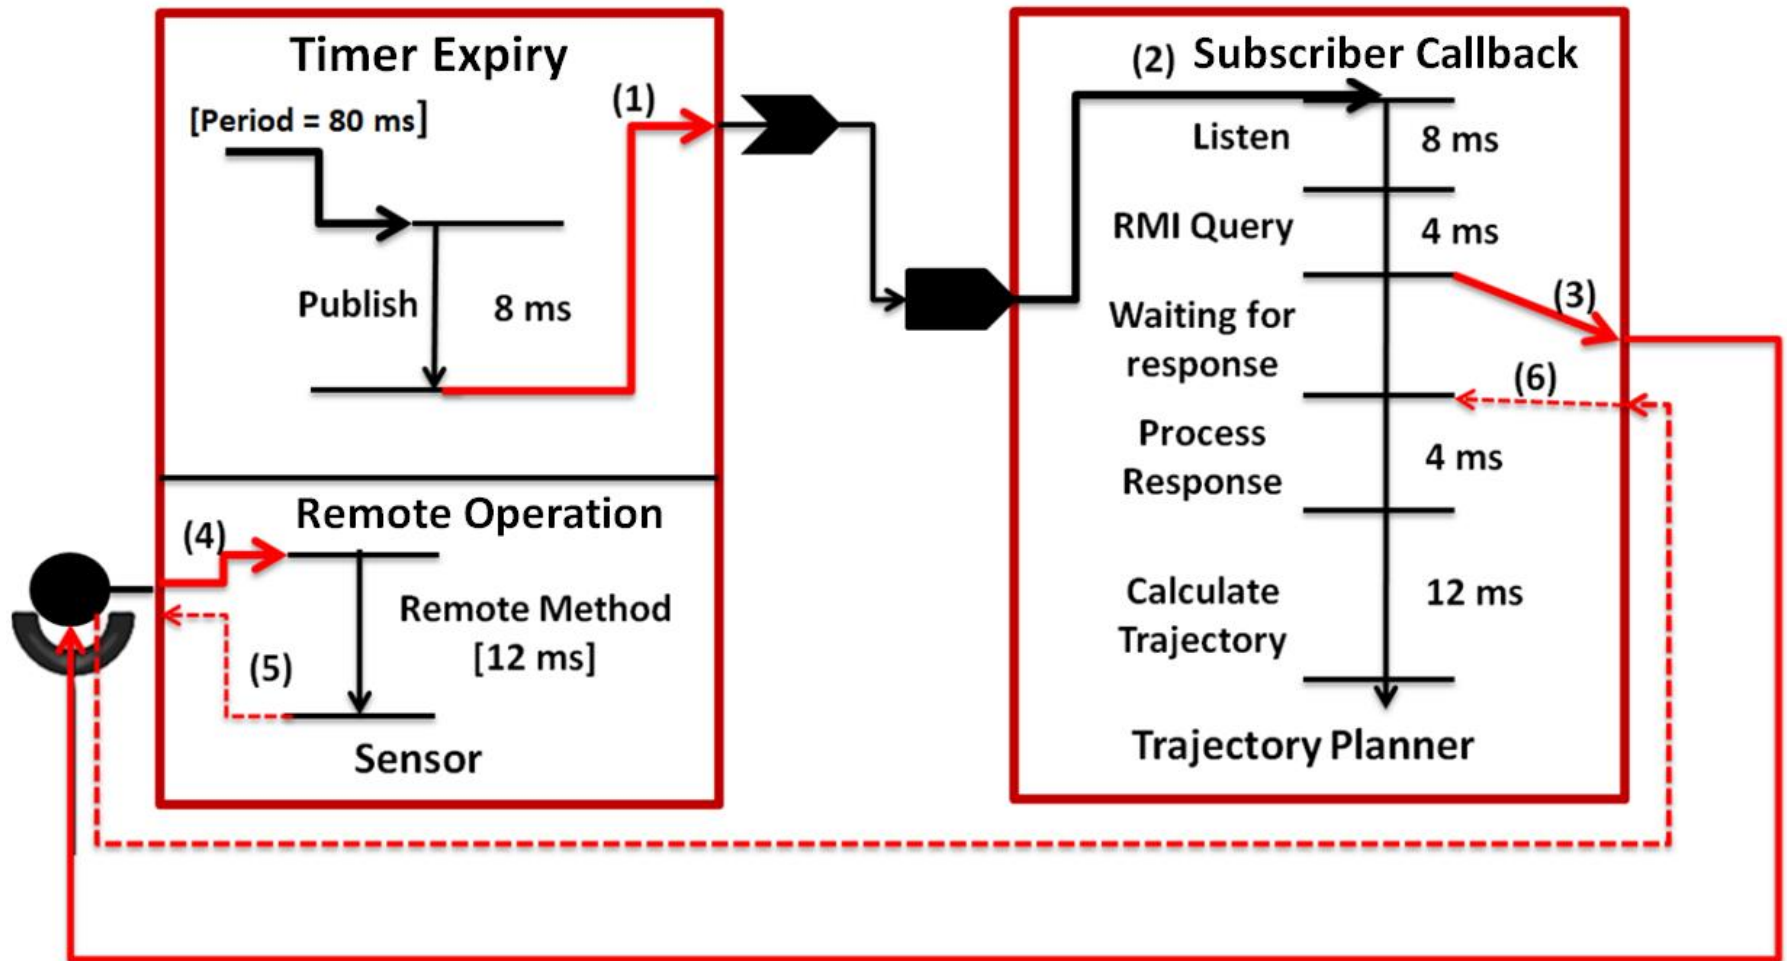
operation_name      =   ID ;
operation_priority  =   INT ;
operation_deadline  =   INT ;
functional_step     =   {sequential_code_block | rmi_call | ami_call | dds_publish | dds_pull_subscribe |
                        dds_push_subscribe | loop} ;

sequential_code_block = INT, ',' ;
rmi call            =   'RMI', ws, receptacle_port, '.', remote_operation, '[' query_time, processing_time ']' ;
ami call            =   'AMI', ws, receptacle_port, '.', remote_operation, '[' query_time, processing_time ']' ;
dds publish         =   'DDS_Publish', ws, dds_port, '.', topic, '[' , publish_time, ']' ;
dds pull subscribe  =   'DDS_Pull_Subscribe', ws, dds_port, '.', topic, '[' , processing_time, ']' ;
dds push subscribe  =   'DDS_Push_Subscribe', ws, dds_port, '.', topic, '[' processing_time, ']' ;
loop                =   'LOOP', ws, '[' , count, ']', ws, '{', {functional_step}, '}' ;
receptacle_port     =   ID ;
remote_operation    =   ID ;
dds_port            =   ID ;
topic               =   ID ;
query_time          =   INT ;
processing_time      =   INT ;
publish_time        =   INT ;
count               =   INT ;
```

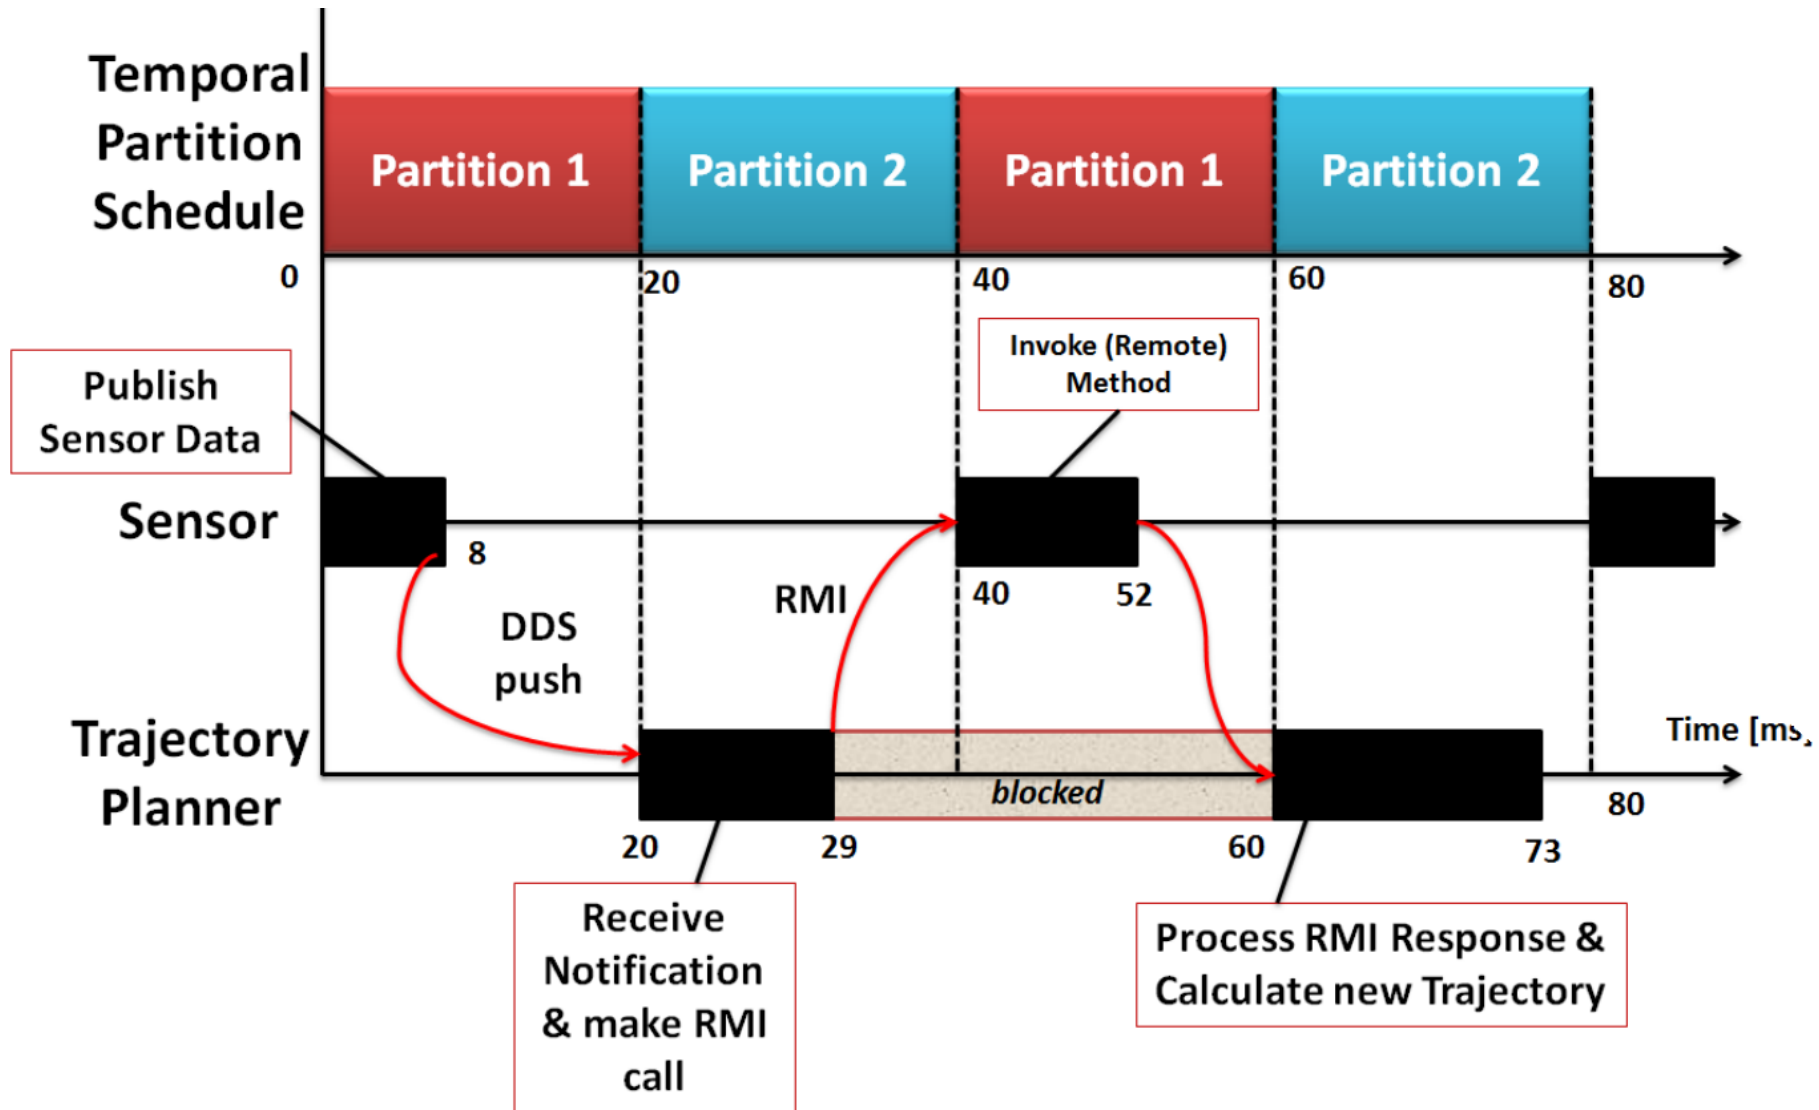
A simple textual language for representing the temporal behavior of a component operation.  
Fixed, worst-case behavior (no data dependency)

---

# Trajectory Planner (1 / 2)



# Trajectory Planner (2/2)



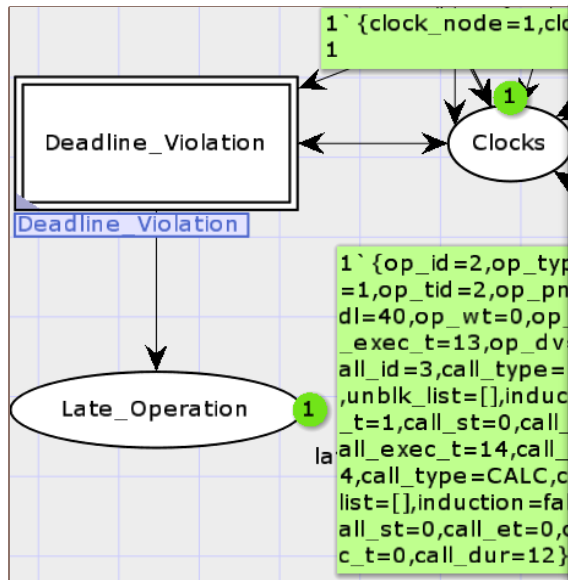


# State Space Analysis

---

- ▶ Generate a bounded state space; a tree of possible executions from the initial state in the CPN model
  - ▶ The analyzable states of this system are observed in the markings of the various CPN places in the model
  - ▶ Using both standard and user-defined queries, the state space is searched to check for system properties
    - ▶ Deadline Violations in component operations
    - ▶ System-wide deadlocks
    - ▶ Worst-case trigger to response times for a known trigger and response operation
    - ▶ Partial thread execution order generation based on timing requirements
-

# State Space Analysis – Deadline Violation



```
val DeadlineViolation = fn : Node -> bool
val Get_Violation_List = fn : Node list -> late_opn ms list
val LateOperation_nodes =
  [99,98,97,96,95,94,93,92,91,90,89,88,87,86,85,84,83,82,81,80,79,78,77,76,75,
   74,73,72,71,70,69,68,67,66,65,64,63,62,61,60,59,58,101,100] : Node list
val sortedLateOperation_nodes =
  [58,59,60,61,62,63,64,65,66,67,68,69,70,71,72,73,74,75,76,77,78,79,80,81,82,
   83,84,85,86,87,88,89,90,91,92,93,94,95,96,97,98,99,100,101]
: CPN'ColorSets.IntCS.cs list
val FirstDeadlineViolation =
  [{op_calls=[{call_blk_t=0,call_dur=14,call_et=0,call_exec_t=14,call_id=3,
    call_op_id=2,call_st=0,call_type=RMI_c,induction=true,pr_t=1,
    q_t=1,unblk_list=[]},
    {call_blk_t=0,call_dur=12,call_et=0,call_exec_t=0,call_id=4,
    call_op_id=2,call_st=0,call_type=CALC,induction=false,pr_t=0,
    q_t=0,unblk_list=[]}],op_dl=40,op_dv=true,op_et=0,
    op_exec_t=13,op_id=2,op_nid=1,op_pn=2,op_prio=5,op_st=20,op_tid=2,
    op_type=DDS_OP,op_wt=0}] : late_opn ms
```

```
fun DeadlineViolation n = (Mark.New_Page'Late_Operation 1 n <> []);
```

```
fun Get_Violation_List [] = []
  | Get_Violation_List (first_node::rest) =
    (Mark.New_Page'Late_Operation 1 first_node)::(Get_Violation_List rest);
```

```
val LateOperation_nodes = SearchNodes (
  EntireGraph,
  fn n => (DeadlineViolation n),
  NoLimit,
  fn n => n,
  [],
  op ::);
```

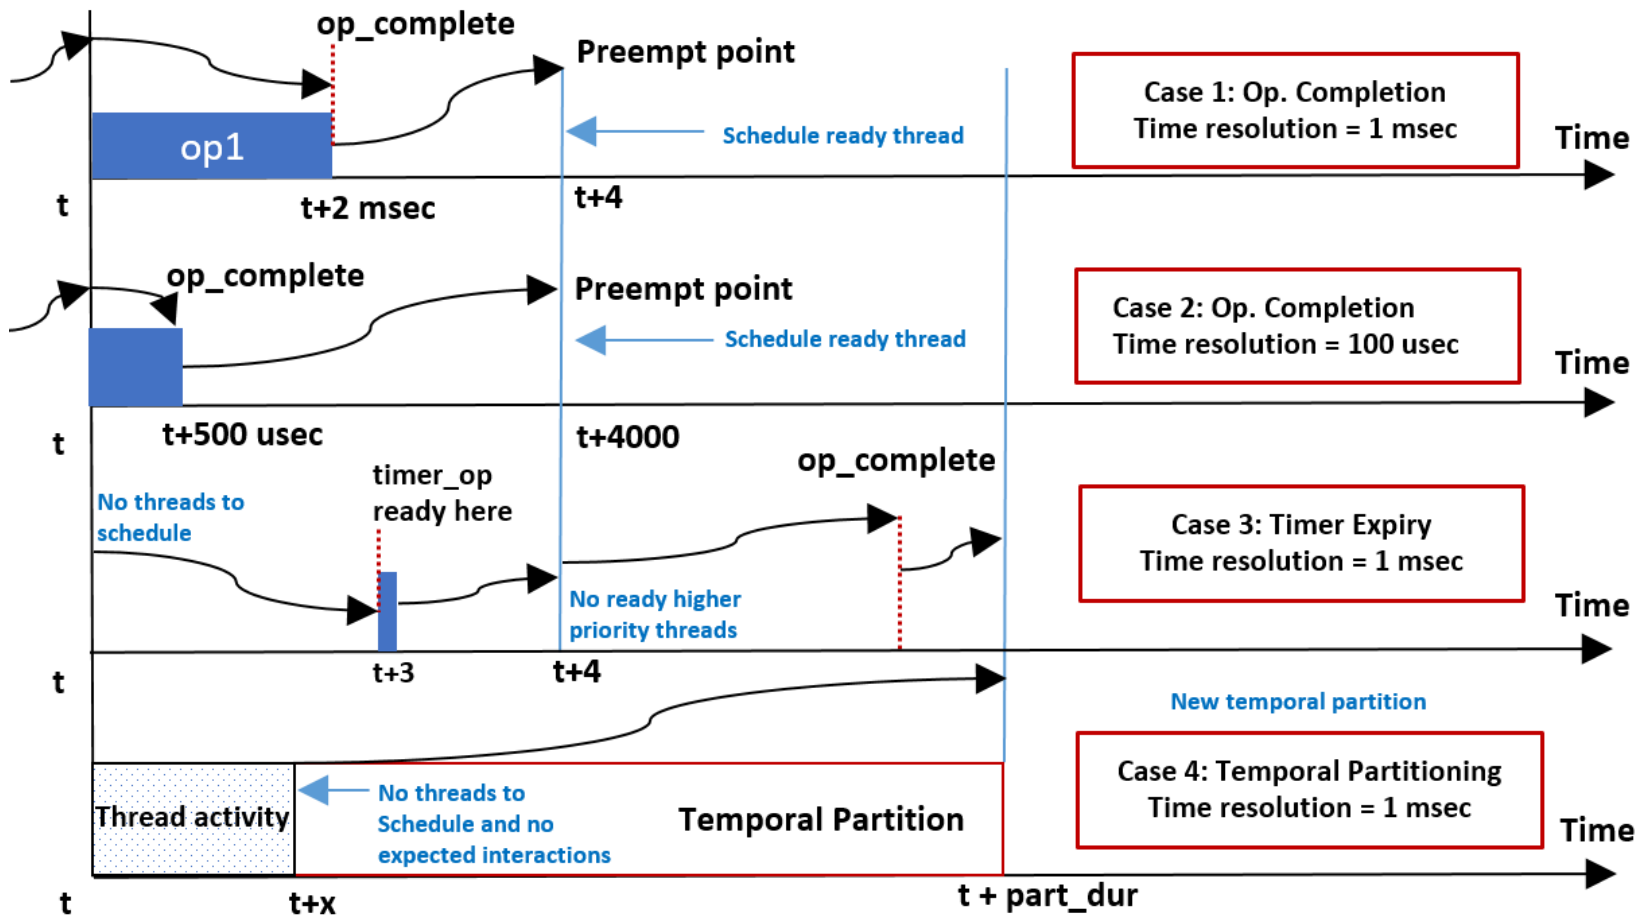
```
val sortedLateOperation_nodes = sort INT.lt LateOperation_nodes;
val FirstDeadlineViolation = (hd (remdupl(Get_Violation_List sortedLateOperation_nodes)));
```

# Analysis Improvements: Handling Time

---

- ▶ Explicit modeling of time as an integer-valued *clock* color token in CPN
  - ▶ Modeling the OS scheduler clock this way allows for easy extensions to its data structure to support
    - ▶ Intermediate time stamps and internal state variables
    - ▶ Adding temporal partitioning and other prioritization schemes
  - ▶ Reduces the total number of colors required by the complete model
  - ▶ Chosen a time quantum is 1 msec (per typical 1KHz scheduler in Linux)
-

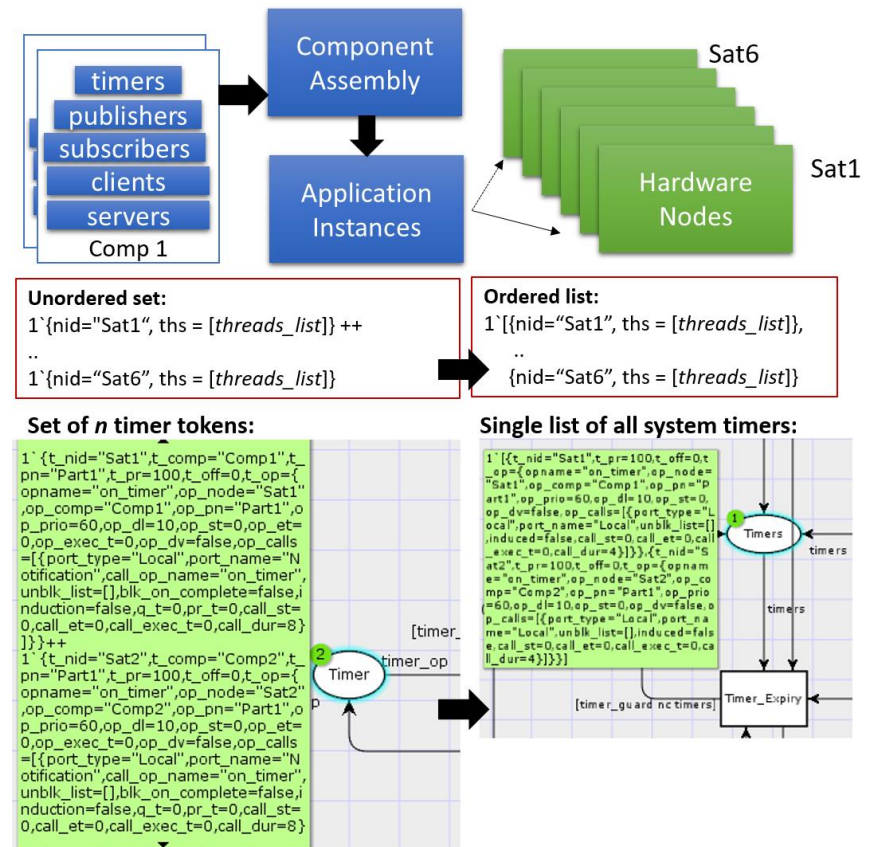
# Smarter Handling of Time Progression



Fast-forward time up to the next relevant event (DEVS approach)

# Analysis Improvements: Distributed Deployment

- ▶ Early models of distributed deployments included a unique token per CPN place for each hardware *node* in the scenario
  - ▶ Lead to non-determinism in transition bindings
- ▶ Employ structural reduction
  - ▶ Merge hardware *node* tokens into a single *list* of tokens instead of an unassociated grouping of node tokens
  - ▶ Simultaneous events happen in all nodes at the same time
- ▶ *Reduces the resultant state space and dramatically improves scalability*



# Advanced State Space Methods

---

- ▶ Compute all reachable states of the modeled system
  - ▶ Derive a directed graph representing the *state space*: the tree of possible executions that the system can take from an initial state
  - ▶ Usage:
    - ▶ Verify behavioral properties such as lack of queue overflows, deadline violations and deadlocks
    - ▶ Derive counter examples if a property is violated
  - ▶ Potential for state space explosion
    - ▶ Needs advanced memory management methods to make state space analysis efficient
-

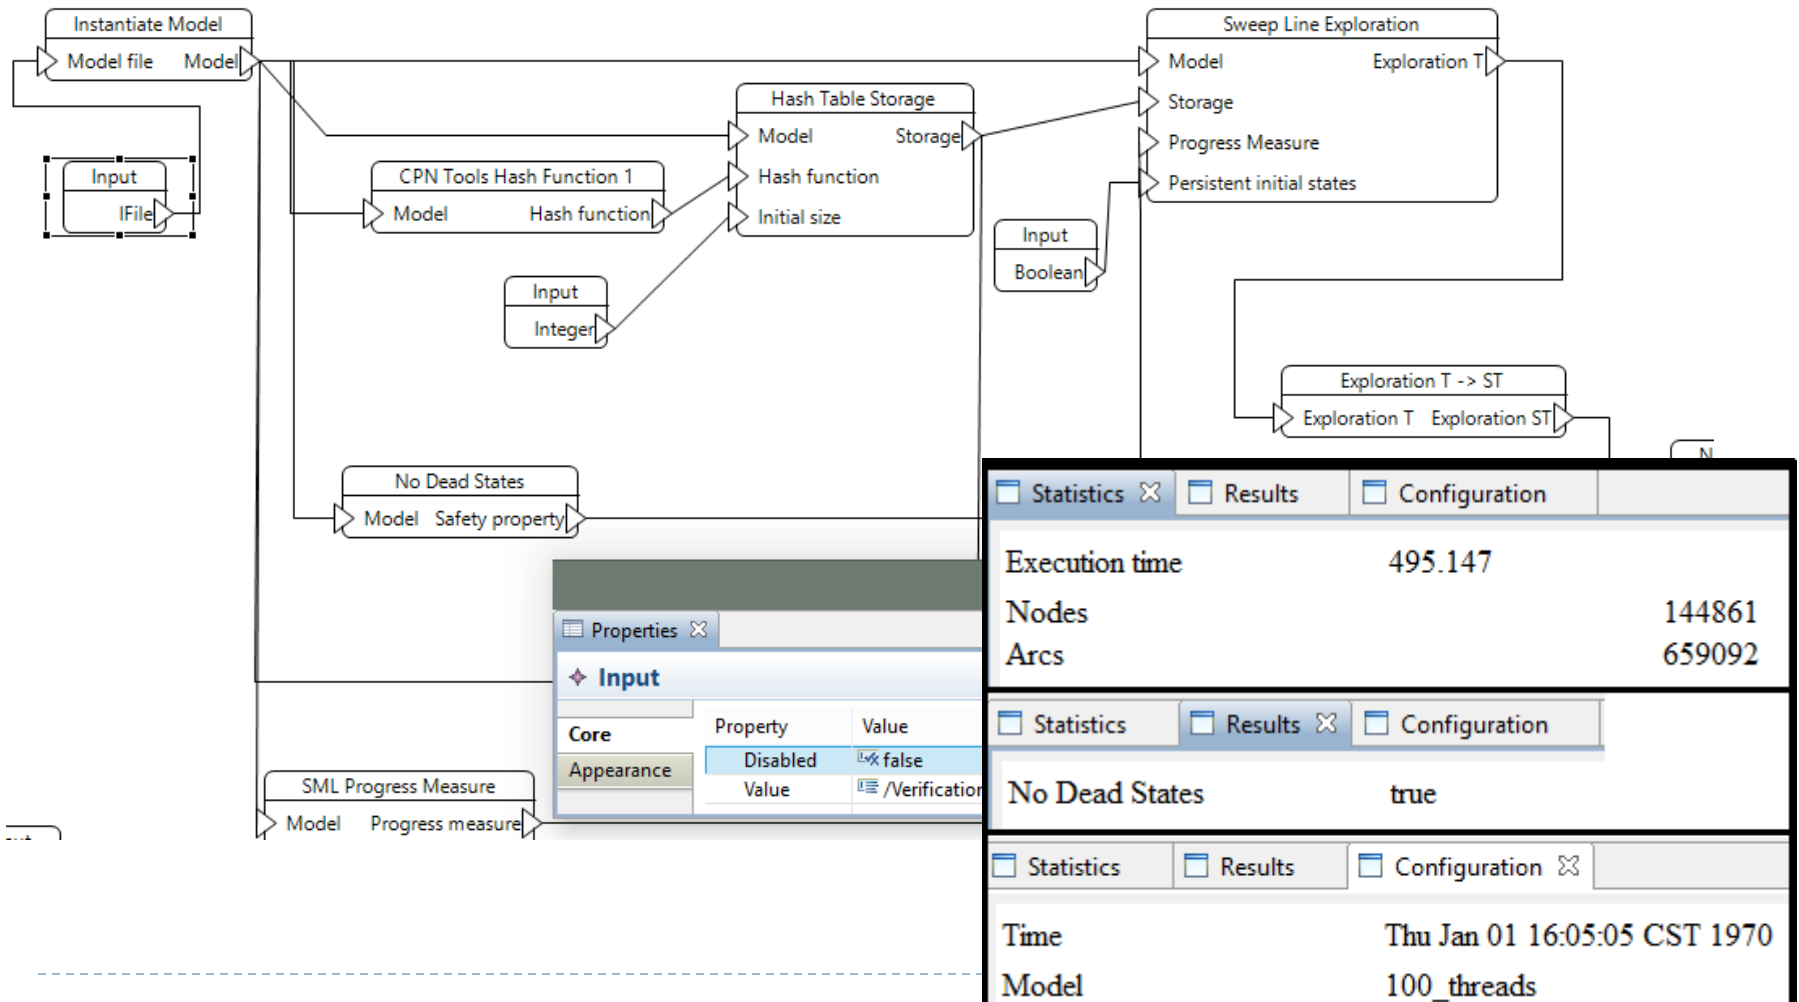
# Advanced State Space Methods

---

- ▶ In order to easily apply advanced state space reduction techniques, we use a tool called **ASAP**
  - ▶ Sweep line method
    - ▶ Discard generated states on-the-fly by performing verification checks during generation time
    - ▶ Any state that does not violate system properties is deleted
  - ▶ **100 interacting components in 10 computing nodes:**
    - ▶ Using CPN Tools built in state space generation:
      - ▶ 20 hyperperiods of activity → 36 minutes on a typical laptop
    - ▶ Using ASAP and on-the-fly verification
      - ▶ Less than 10 minutes to perform deadlock checks on the deployed system
-

# Advanced State Space Methods

ASAP screenshot:





# Future Work

---

- ▶ DREMS component communication is facilitated by a time-varying network
    - ▶ Bandwidth provided by the system predictably fluctuates between a minimum and maximum (e.g. due to orbital period)
    - ▶ Currently we assume worst-case network delay
    - ▶ Work in progress: capturing the *network profile* (network performance over time) of a deployment
  - ▶ Investigating the utility of this approach on fault-tolerant and self-adaptive systems
    - ▶ Integration with a run-time resilience engine
    - ▶ Checking for timing anomalies before settling on a reconfiguration strategy
-

# Conclusions

---

- ▶ DREs running time-critical applications must satisfy strict timing requirements to operating safely
  - ▶ To reduce design and integration complexity, component-based design models are increasingly being used
  - ▶ Appropriate analysis models are required to study the structural and behavioral complexity of such designs
  - ▶ Model-based development tools integrated with analysis tools and code generation offer an integrated solution: what is analyzed is the same what runs in the executing system
  - ▶ Access:
    - ▶ <https://drems.isis.vanderbilt.edu> – complete system
    - ▶ Github: rosmo project – next generation based on ROS
-

---

Thank You

---

# Overview and Outline

---

- ▶ **Distributed Real-time Managed Cyber-Physical Platforms**
    - ▶ Distributed Applications
    - ▶ Challenges
  - ▶ Layered architecture for building a distributed software platform for these systems
  - ▶ Overview of the DREMS platform, a prototype implementation.
  - ▶ Example
-

# Distributed Real-time Managed Cyber-Physical Platforms 1/2

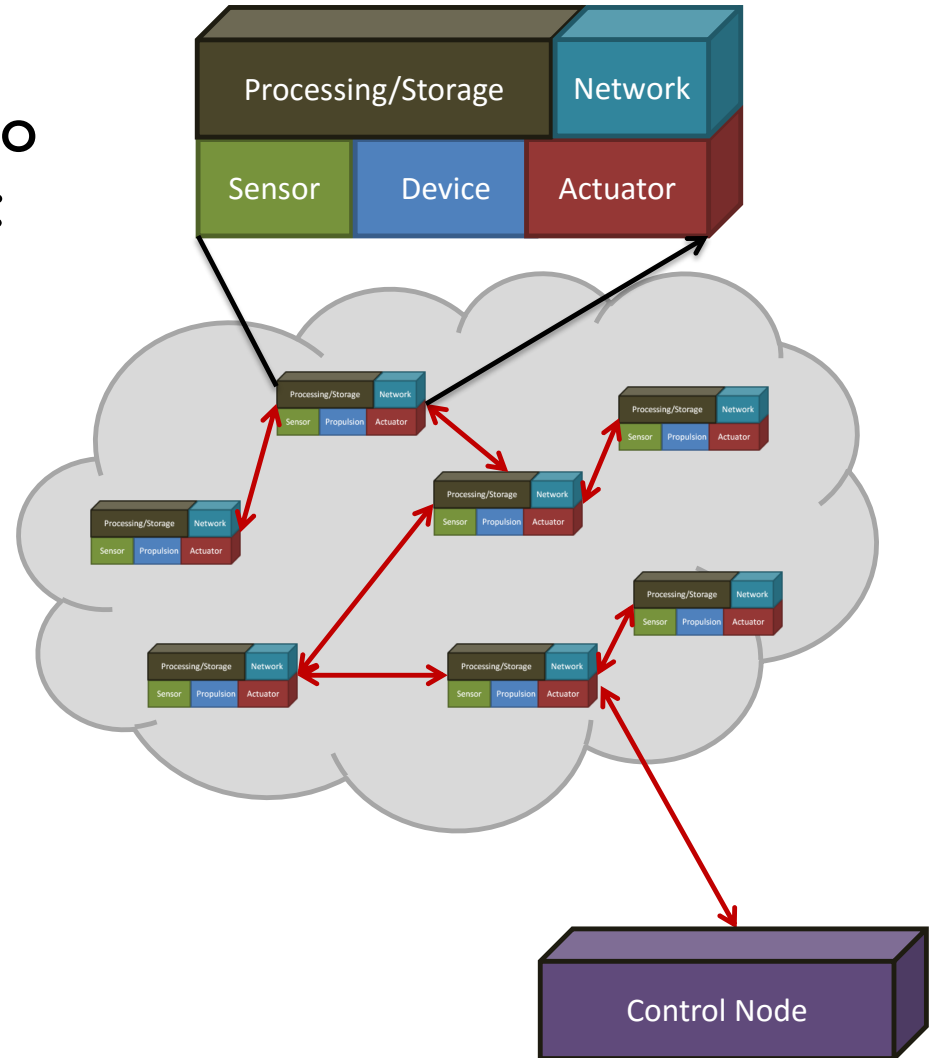
---

- ▶ Built not as a single use, single function network, but as networked (wireless) platforms that can be used by many, possibly concurrent users
  - ▶ Physical configuration/topology affects the available computational resources.
  - ▶ Physics imposes timing constraints on the computational and communication activities.
  - ▶ Critical system software is required to verifiably meet the design requirements.
- 



# Distributed Real-time Managed Cyber-Physical Platforms 2/2

- ▶ Applications span multiple nodes, for reasons related to the availability of resources:
  - ▶ some nodes may have sensors,
  - ▶ some may have actuators,
  - ▶ some may have computing,
  - ▶ some may have storage resources.
- ▶ Applications must be architected to rely on loosely interacting components.



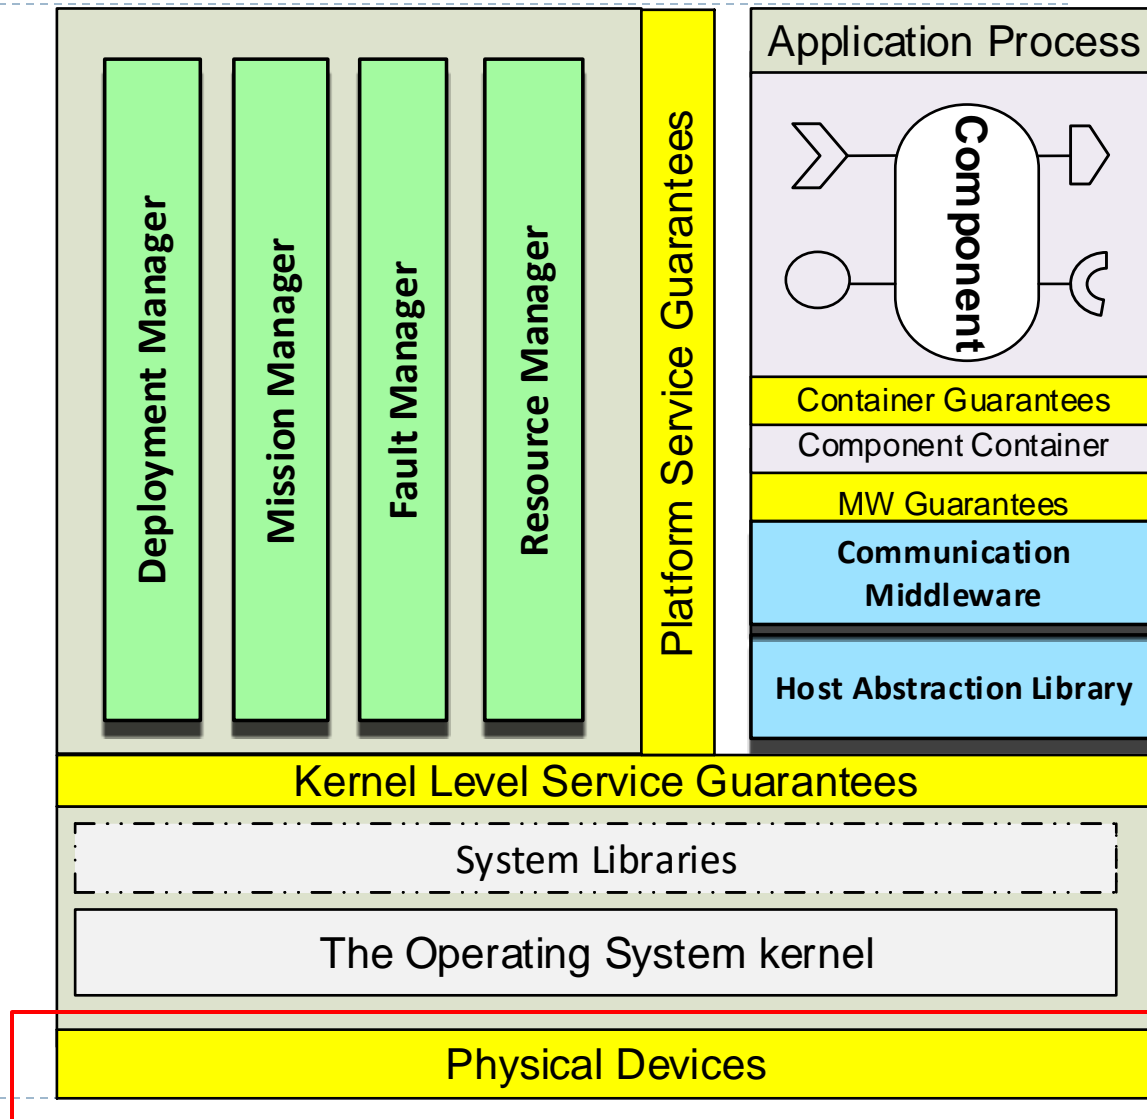
# Challenges

---

- ▶ Remote deployment and configuration
  - ▶ Resilient operation requirements
  - ▶ Share applications from different vendors and users with different privileges.
  - ▶ Information sharing/ leakage between applications must be controlled under an overall system security policy
  - ▶ Performance isolation is critical
    - ▶ One application should not be able to affect the functionality or performance of another application
    - ▶ The absence of strong performance isolations will also result in security problems
-

# Layered architecture approach

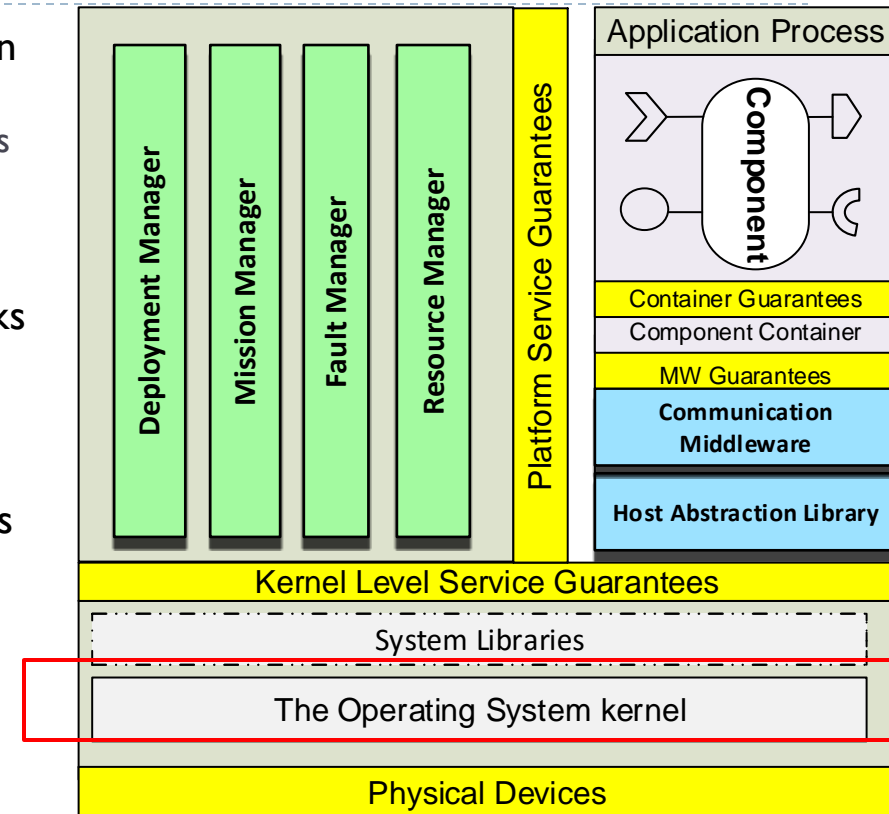
- ▶ Watchdog timers reset the platform upon deadlock or critical failures.
- ▶ Robust Networking
  - ▶ support time-constrained and real-time communications with guarantees
  - ▶ Provides updates on channel bandwidth and expected latency map.
  - ▶ support multiple traffic classes natively
  - ▶ Provide protection against external interference
  - ▶ Support link level encryption





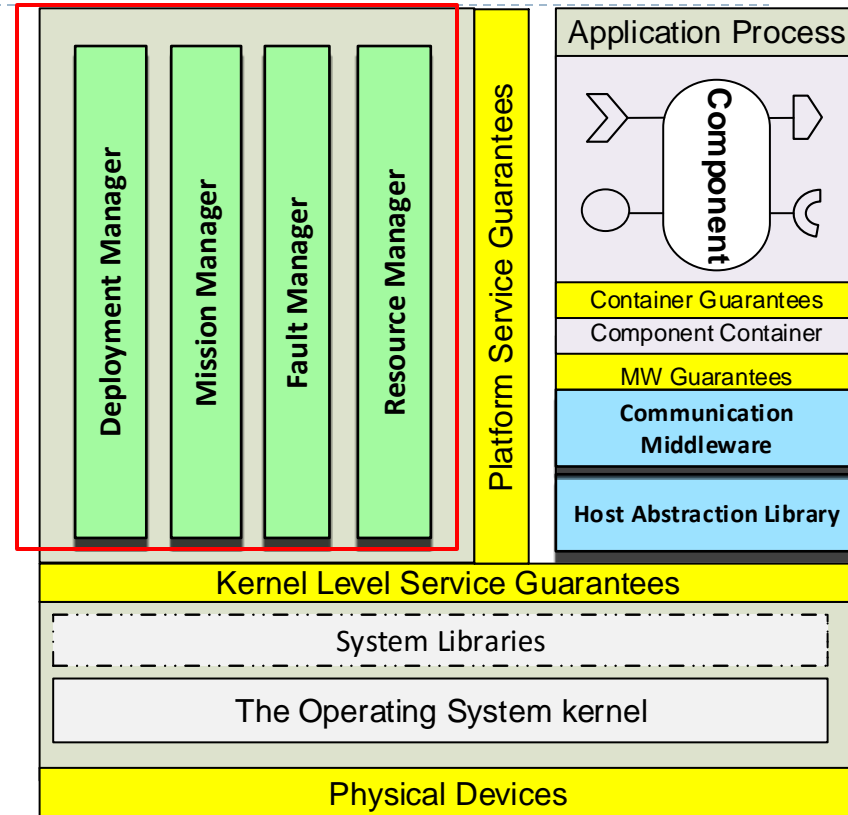
# Trusted Computing Base: Operating System

- ▶ Support for multiple critical levels of computation tasks
  - ▶ Scheduling policies that provide verifiable guarantees for the timeliness of task execution.
  - ▶ Strict isolation between applications of different criticality levels.
- ▶ Multiple critical levels of communication networks
  - ▶ Different traffic classes.
  - ▶ Reduction in covert channels of communication.
  - ▶ Real-time support for communication.
- ▶ Fault tolerant clock synchronization across nodes over wireless network is required
- ▶ Confidentiality, Integrity and Authenticity guarantees for communication (require cryptography devices)
  - ▶ Mandatory access control and multi-level security may be required.
- ▶ Rich and extensible capability model
  - ▶ Control what platform services can be used by an application
  - ▶ Support for easy migration of application processes between nodes without affecting other applications



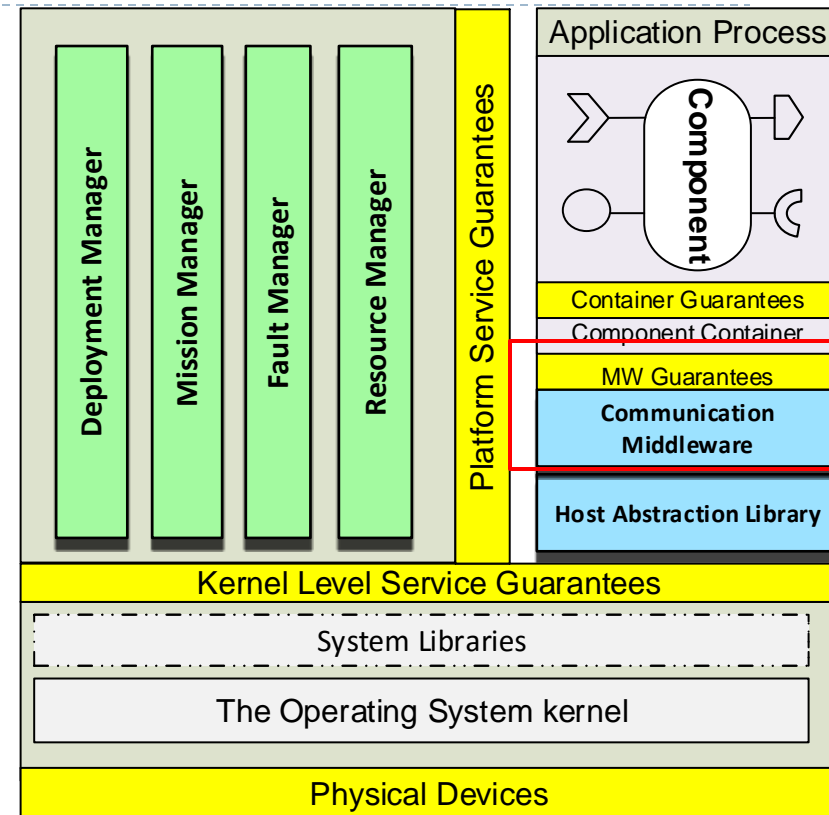
# Trusted Computing Base: Platform Services

- ▶ Privileged runtime software actors are required for system level tasks
- ▶ Deployment Management
  - ▶ Fault tolerant application deployment and configuration.
- ▶ Mission management
  - ▶ Activate/Deactivate applications based on events/passage of time.
- ▶ Fault Management
  - ▶ Restore mission/system functionality without external intervention
- ▶ Resource Management
  - ▶ implement a dynamic resource allocation policy, where applications can dynamically request and release resources, and the service honors or rejects these requests while maximizing system utility



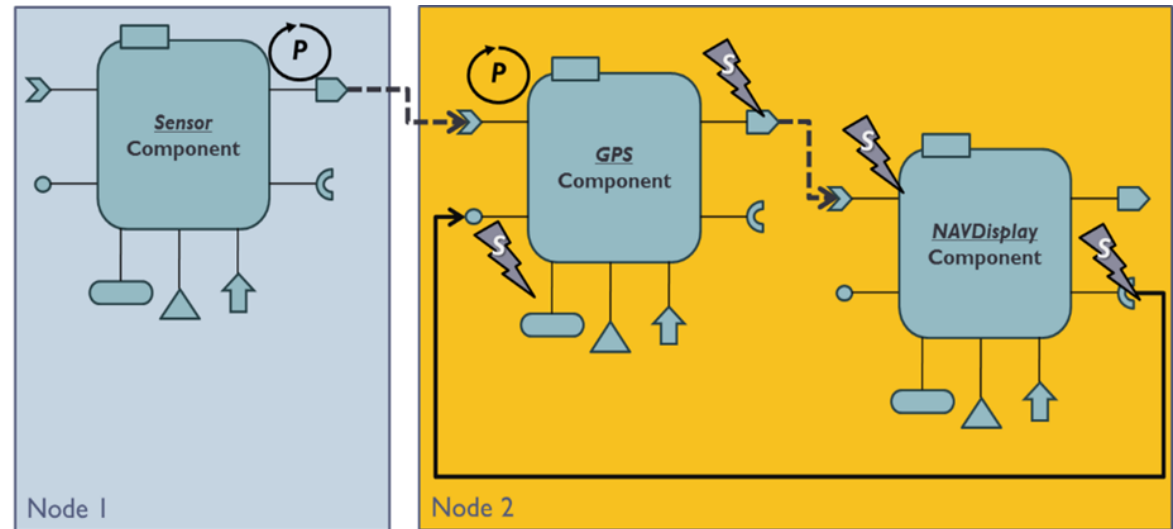
# Middleware

- ▶ Well defined interaction patterns, e.g.
  - ▶ Point to point
  - ▶ Pub/Sub
- ▶ Provide abstractions to configure and use the different traffic classes and the quality of service information provided by the networking layer
- ▶ Support for store and forward for non-real-time communication flows.
- ▶ Support for multiple levels of security at the middleware level.




# Robust Software Component Model and Development Tools

- ▶ Clearly delineate computational aspect from communication aspect.
- ▶ Precise scheduling model for the operations.



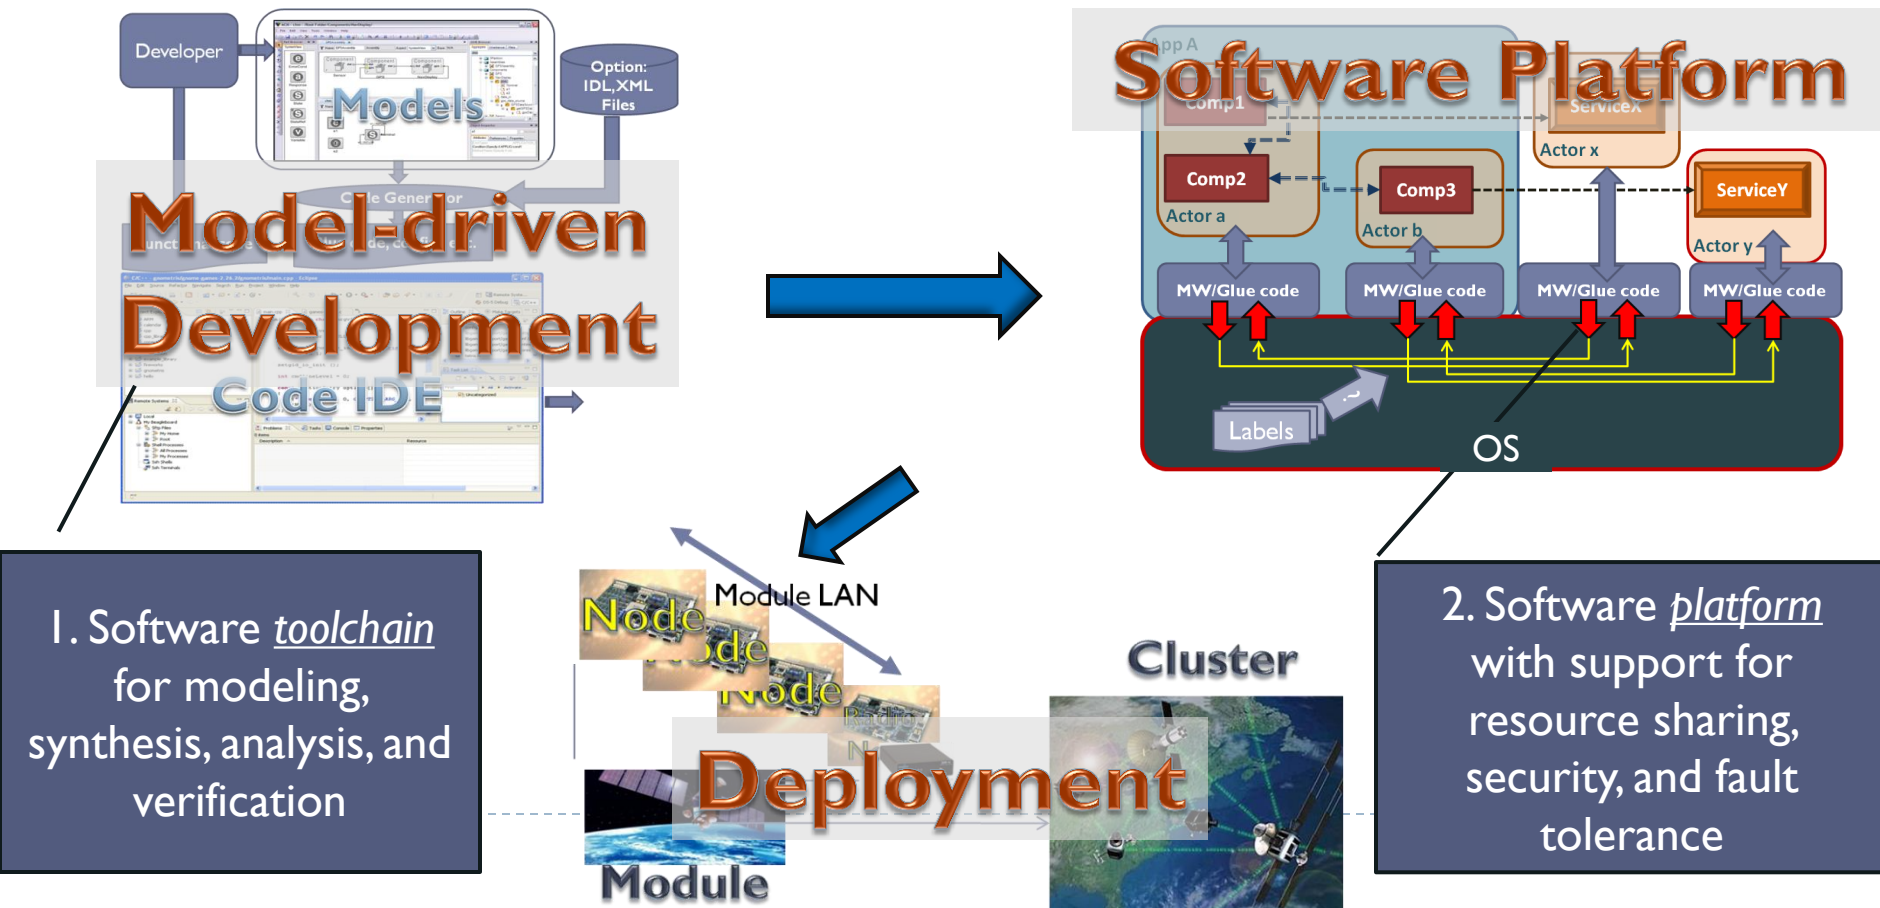
- ▶ Support for security policies. Should also provide fault management, including anomaly detection, diagnosis and fault mitigation.
- ▶ Prevent technology lock-in:
  - ▶ Support for different programming languages
  - ▶ Should not require any particular middleware implementation.
- ▶ Support for integrated development tools that reduce the time to develop and integrate new applications into the system.



## Preliminary Results: DREMS toolchain and platform

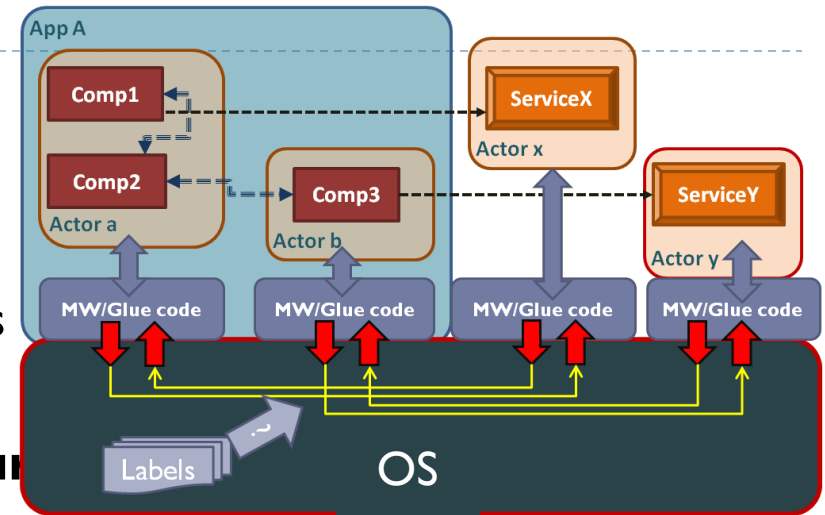
# DREMS

- Distributed REaltime Managed System. DREMS consists of:
1. A software *platform*, consisting of an OS and middleware
  2. A software *toolchain*, for modeling applications



# Operating System

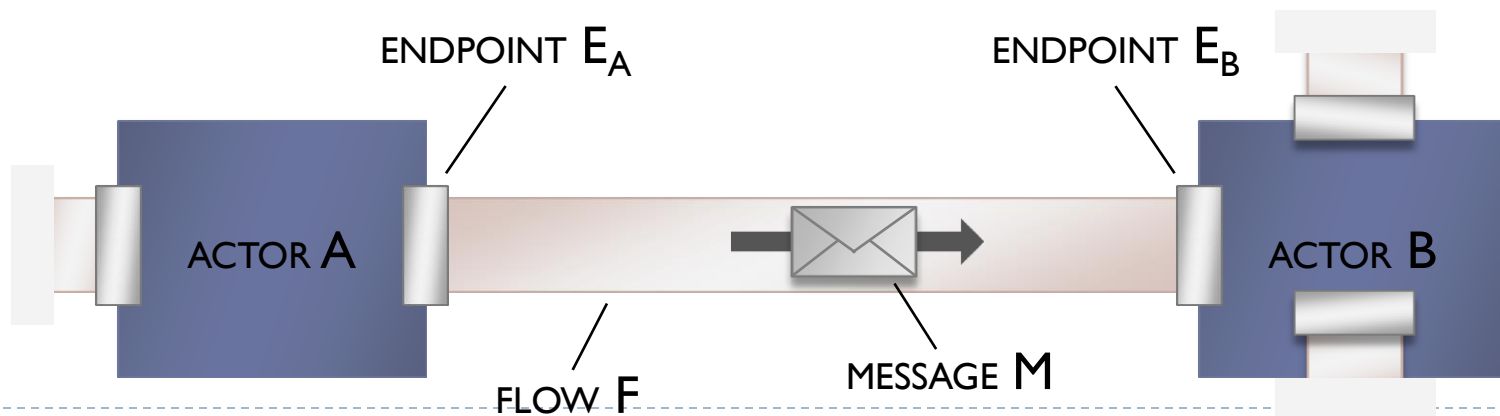
- ▶ Supports actor concept.
- ▶ Resource sharing is strictly monitored and controlled.
- ▶ Fine-grained capability model that controls access to different system services.
- ▶ All interactions among actors are via '**secure transport**'.
- ▶ Part of the Trusted Computing Base that enforces the MLS/MAC security policies.
- ▶ Real-time scheduling model that supports mixed-criticality systems.
  - ▶ System (platform Actors – uses simple priority scheduling)
  - ▶ Application (uses temporal partitioning)
  - ▶ Best effort (uses CFS)
- ▶ Temporal partitioning. Each partition is temporally separated from others.



# Secure Transport

---

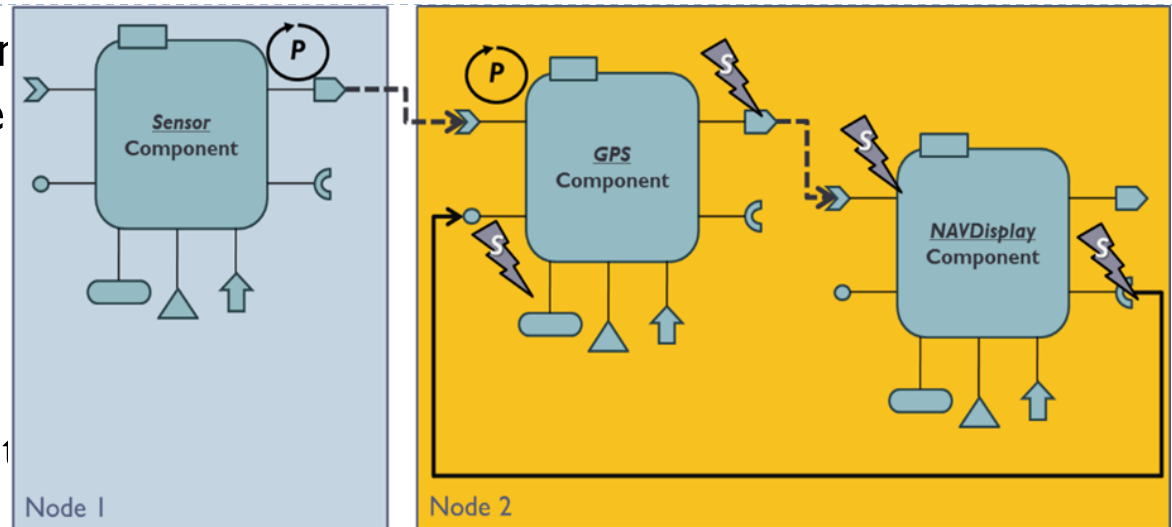
- ▶ The secure transport mechanism enforces the following:
  - ▶ Actors write messages only to **endpoints**.
  - ▶ Endpoints and **flows** are configured only by trusted platform Actors; used by regular Actors.
  - ▶ Enforces mandatory access control for the messages.
  - ▶ Supports both UDP and SCTP protocols.
  - ▶ All messages must have an label assigned *by the sending actor*.
  - ▶ The messages can go to a destination if and only if the destination has a label that **dominates** the label of destination.





# Component-Based Distributed Applications

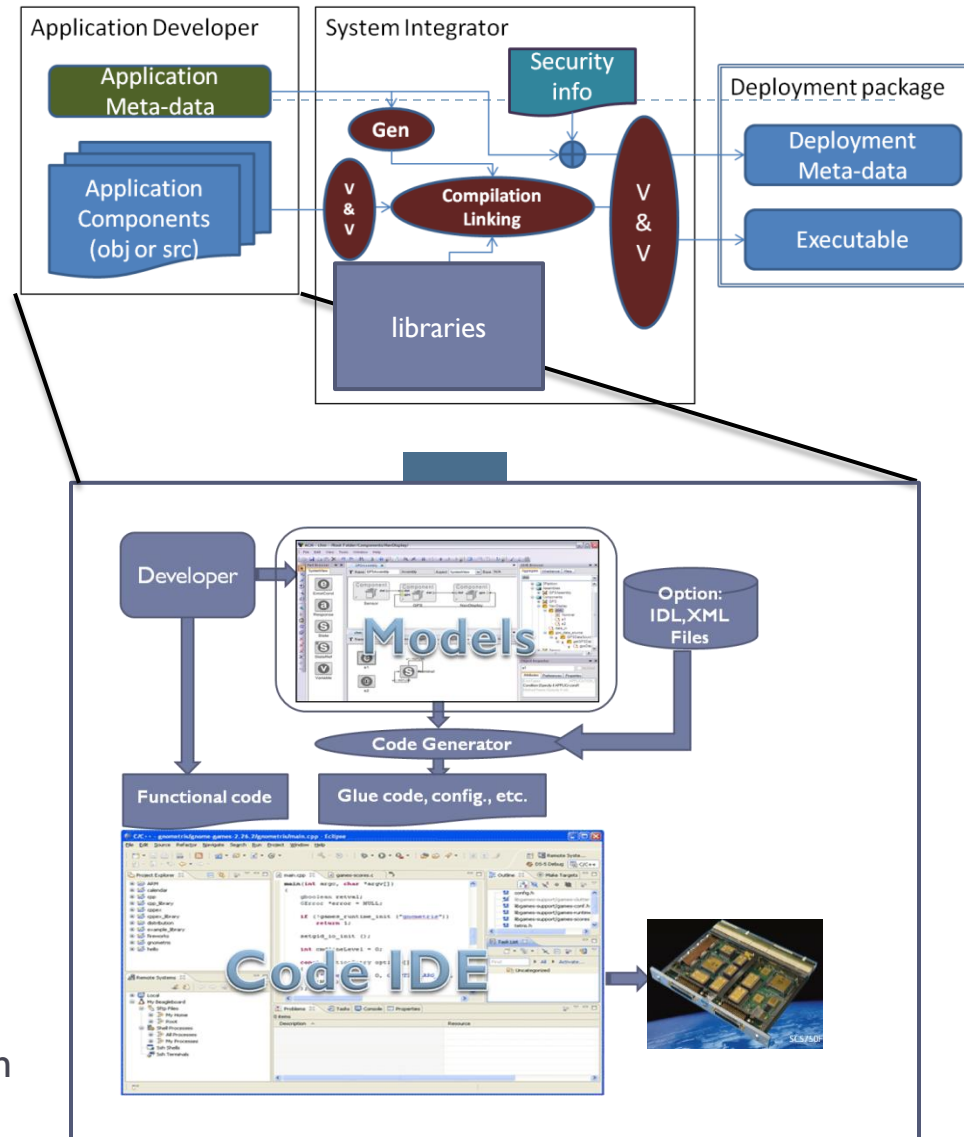
- ▶ Apps are architected as one or more actors that share information via secure messaging.
- ▶ Actors encapsulate components, with well-defined interaction semantics.



- ▶ Clearly delineate computational aspect from communication aspect.
- ▶ Utilizes connectors for adapting to different middleware implementations.
- ▶ Component work is divided into operations which are scheduled one at a time.
- ▶ The business logic code written by developers is free from any synchronization code, which is one of the common mistakes made by developers.
- ▶ Provides different scheduling policies for operations: First-In First-Out, Priority, Earliest Deadline First.

# Development and Integration Environment

- ▶ Designer creates architectural models for actors and components, and their interactions
- ▶ Models are annotated with resource needs, names for security labels, pre/post-conditions, invariants, etc.
- ▶ Software generator tools produce skeleton/glue code for the application modules
  - ▶ Option: import code generated by code generators, like Matlab/Real-time Workshop
- ▶ System integrator verifies and assembles complete software suite
  - ▶ Check information flows for policy violations
  - ▶ Perform admittance tests by checking application communication requirements against expected network resource profile





Example

# Example: DREMS Orbital Satellite Platform 1/2

## Software Bundle on each satellite



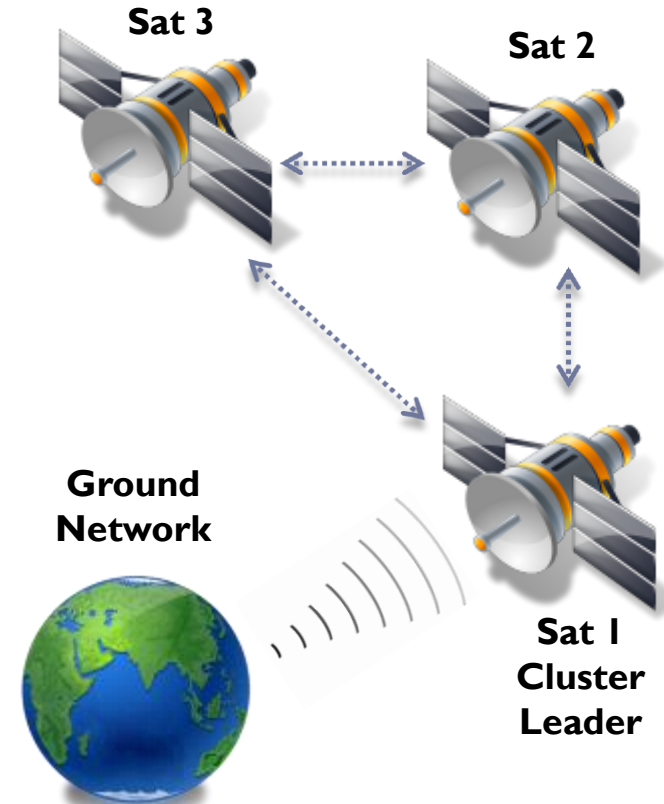
CFA



IPA

Application	Software Components	Partition
Cluster Flight Application (CFA)	Orbital Maintenance	Temporal Partition 1
	Trajectory Planning	System
	Module Proxy	System
	Command Proxy	System
Image Processing Application (IPA)	Image Processing Component (4 on each)	Temporal Partitions 2 & 3

**DESIGN-TIME  
MODELING**



- **Platform management actors** on each satellite provide the capability to **configure, deploy** and **manage** the applications and their actors.
- These actors seamlessly share the CPU with the deployed applications.

**RUN-TIME  
DEPLOYMENT/MANAGEMENT**

# Example: DREMS Orbital Satellite Platform 2/2

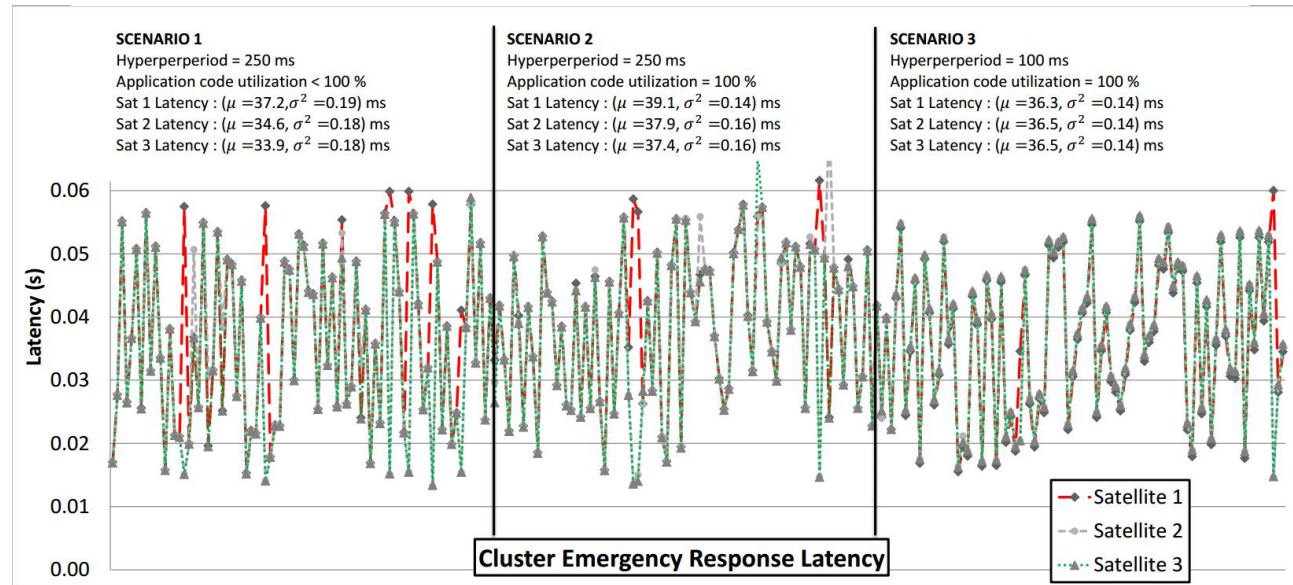
- ▶ The applications were deployed on the platform.
- ▶ Demo integrates the physics model.
- ▶ It shows how a critical flight application can be assembled from components using the development tools.
- ▶ It shows that the image processing application does not affect the CFA response time.
- ▶ The demo shows that the critical flight application is not affected image processing application



BEFORE ENGINE THRUST ACTIVATION  
at time  $t$



AFTER ENGINE THRUST ACTIVATION  
at time  $t + 5$  (min)



# Summary

---

- ▶ Distributed and managed cyber-physical systems pose requirements that go further than traditional DRE systems.
    - ▶ Security, performance isolation, and loose coupling among distributed application are key requirements.
    - ▶ The physics of the platform affect the availability of both computation and communication resources
    - ▶ A layered architecture that builds upon the guarantees provided by the layer below was described.
  - ▶ A prototype platform called DREMS was also discussed.
  - ▶ Please visit: <http://www.isis.vanderbilt.edu/DREMS> for more details.
-